

# Bring Your Own Model

Machine Learning Immersion Day  
Module 3

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



# Agenda

09:00 Welcome and Introductions

09:15 Data Science Lifecycle

10:30 Lab 1: Understanding and preparing data with S3, Glue and Athena

11:15 Lunch

12:00 Model training, testing and deploying with Sagemaker

12:45 Lab 2: Train, test and deploy your first model with Sagemaker

13:45 Break

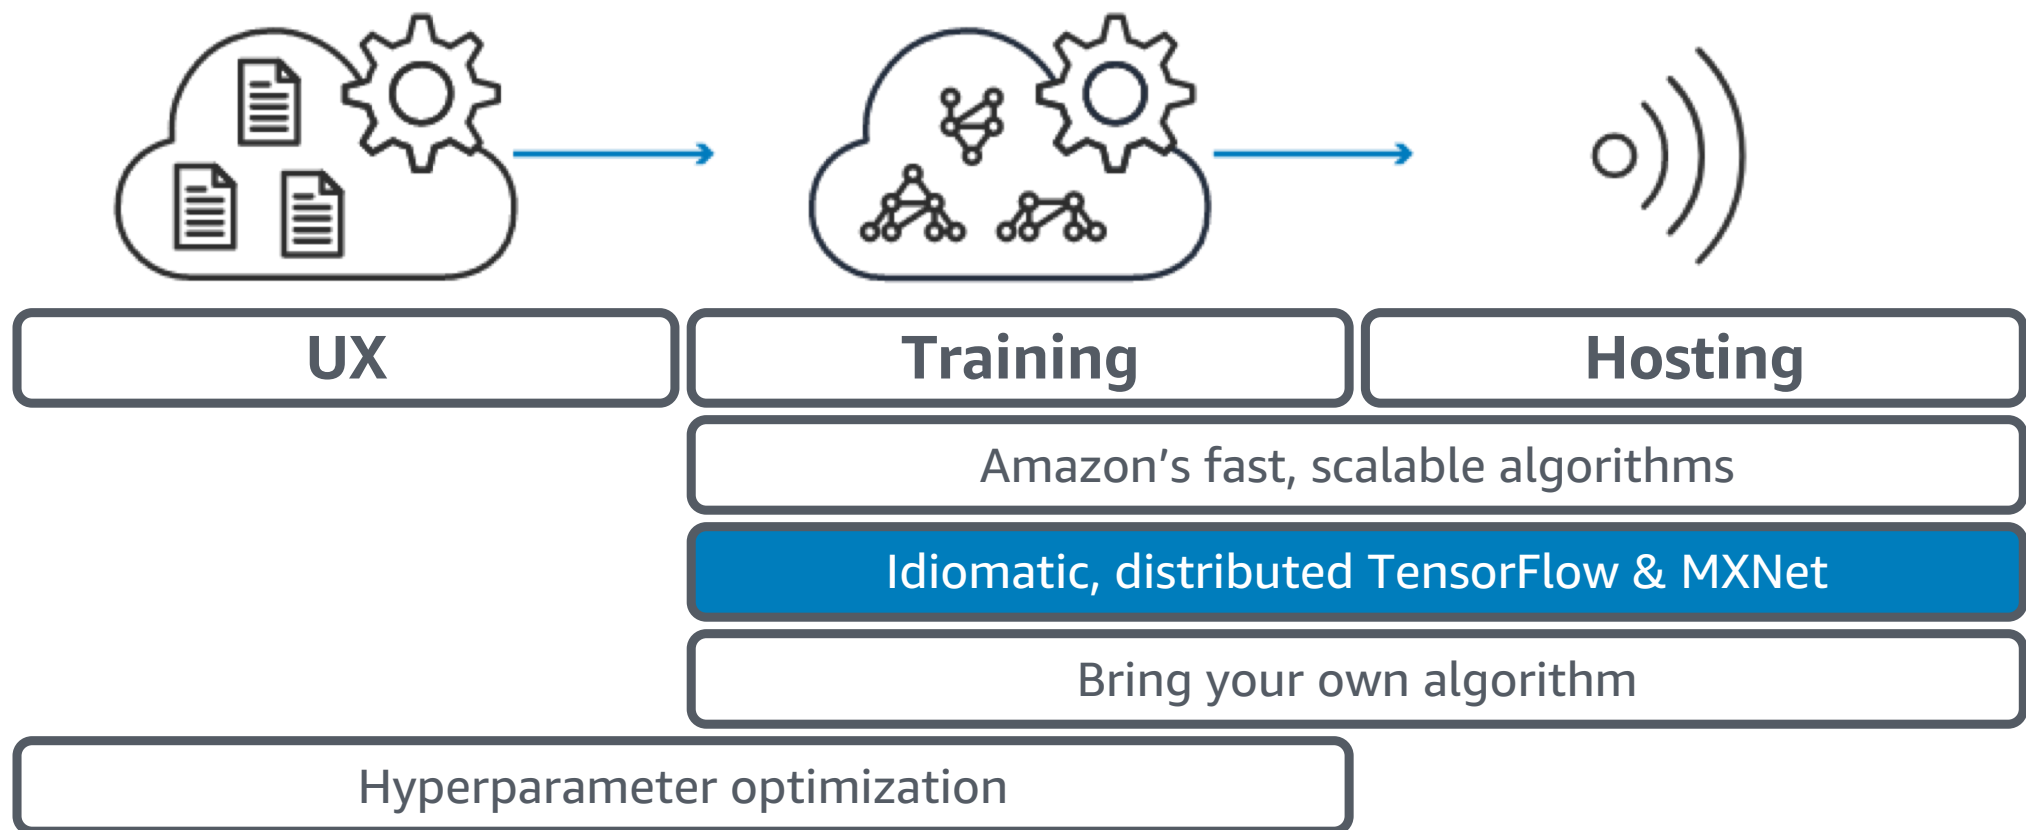
14:00 Continuous Delivery of ML models

14:30 Lab 3: Continuous Delivery of ML models to Amazon SageMaker

**15:15 Bring your own model**

15:45 Wrap Up

# Amazon SageMaker components - Frameworks




# Frameworks



Sample your data...



... explore and refine models in a single Notebook Instance



Use the same code to train on the full dataset in a cluster of GPU instances...



... deploy to production

# Tensorflow example

Use `sagemaker.tensorflow` package to invoke your custom Tensorflow code, provided in a script containing specific functions:

- Exactly one of the following:
  - `model_fn`: defines the model that will be trained.
  - `keras_model_fn`: defines the `tf.keras` model that will be trained.
  - `estimator_fn`: defines the `tf.estimator.Estimator` that will train the model.
- `train_input_fn`: preprocess and load training data.
- `eval_input_fn`: preprocess and load evaluation data.

# Tensorflow example

```
from sagemaker.tensorflow import TensorFlow
tf_estimator = TensorFlow(entry_point='tf-train.py', role='SageMakerRole',
                          training_steps=10000, evaluation_steps=100,
                          train_instance_count=1, train_instance_type='ml.p2.xlarge')
tf_estimator.fit('s3://bucket/path/to/training/data')

# -----

def model_fn(features, labels, mode, hyperparameters):
    # Logic to do the following:
    # 1. Configure the model via TensorFlow operations
    # 2. Define the loss function for training/evaluation
    # 3. Define the training operation/optimizer
    # 4. Generate predictions
    # 5. Return predictions/loss/train_op/eval_metric_ops in EstimatorSpec object
    return EstimatorSpec(mode, predictions, loss, train_op, eval_metric_ops)

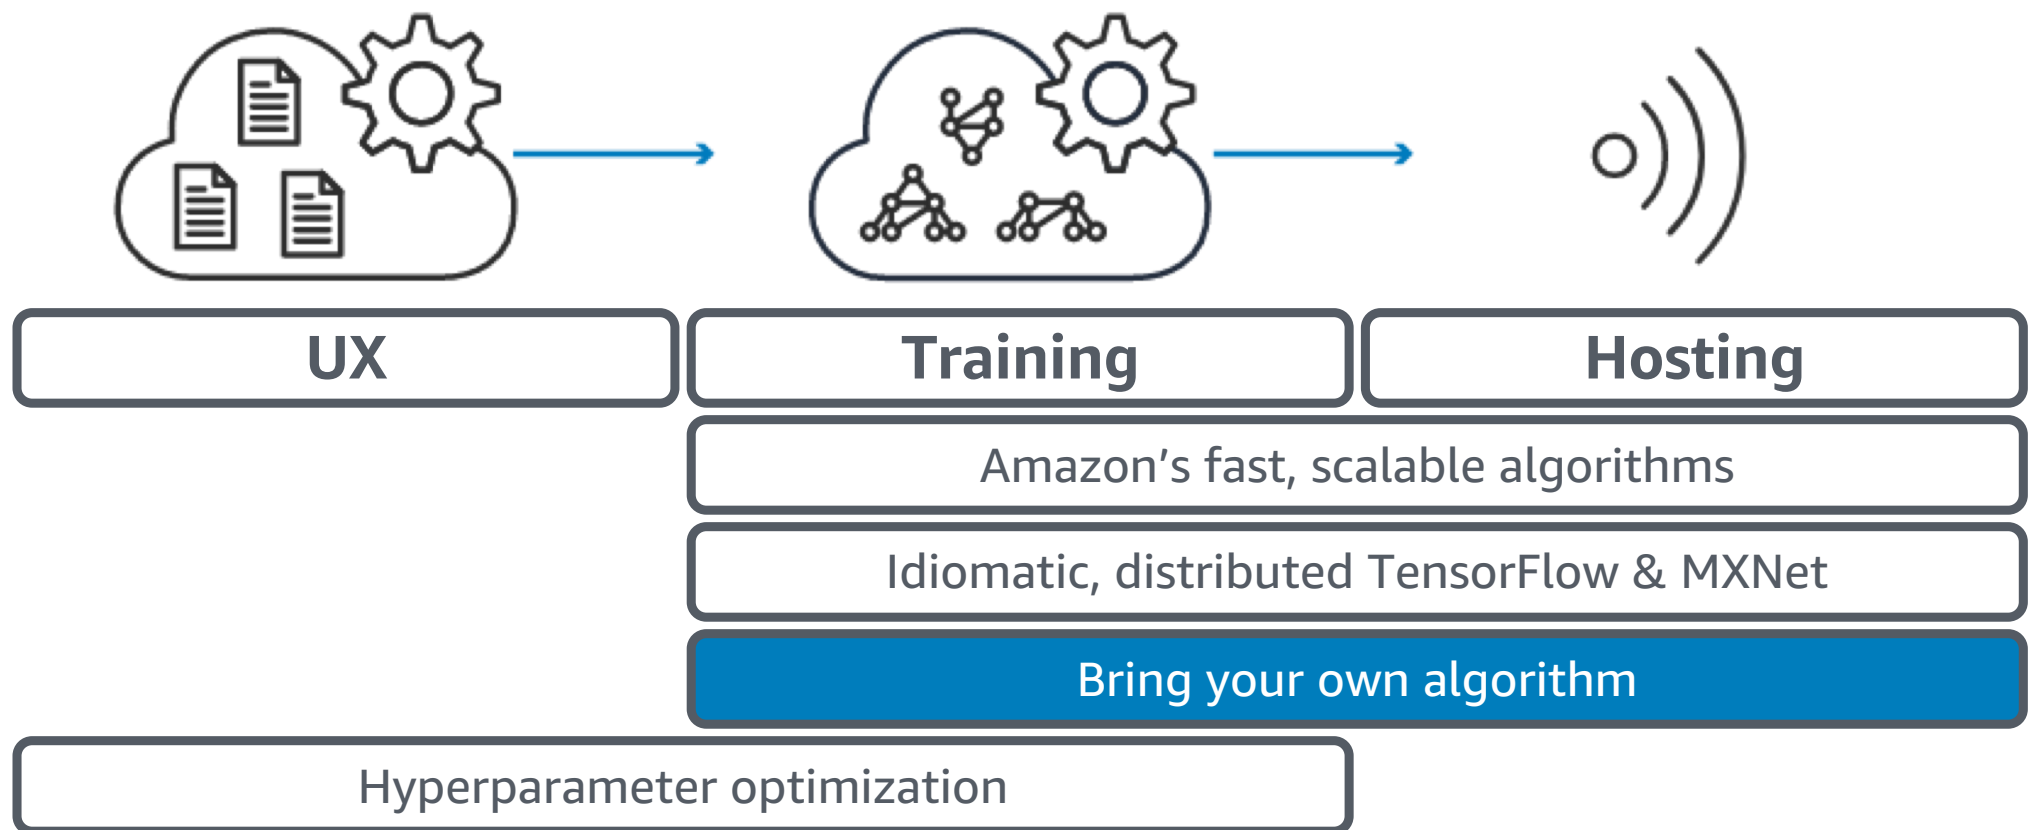
# -----

def train_input_fn(training_dir, hyperparameters):
    # Logic to do the following:
    # 1. Reads the **training** dataset files located in training_dir
    # 2. Preprocess the dataset
    # 3. Return 1) a dict of feature names to Tensors with
    # the corresponding feature data, and 2) a Tensor containing labels
    return features, labels
```


© 2018, Amazon Web Services, Inc. or its Affiliates. All rights reserved.




# Amazon SageMaker components - BYOM



# Bring Your Own Model



Pick your  
preferred  
framework...



... add algorithm  
code to a Docker  
container...

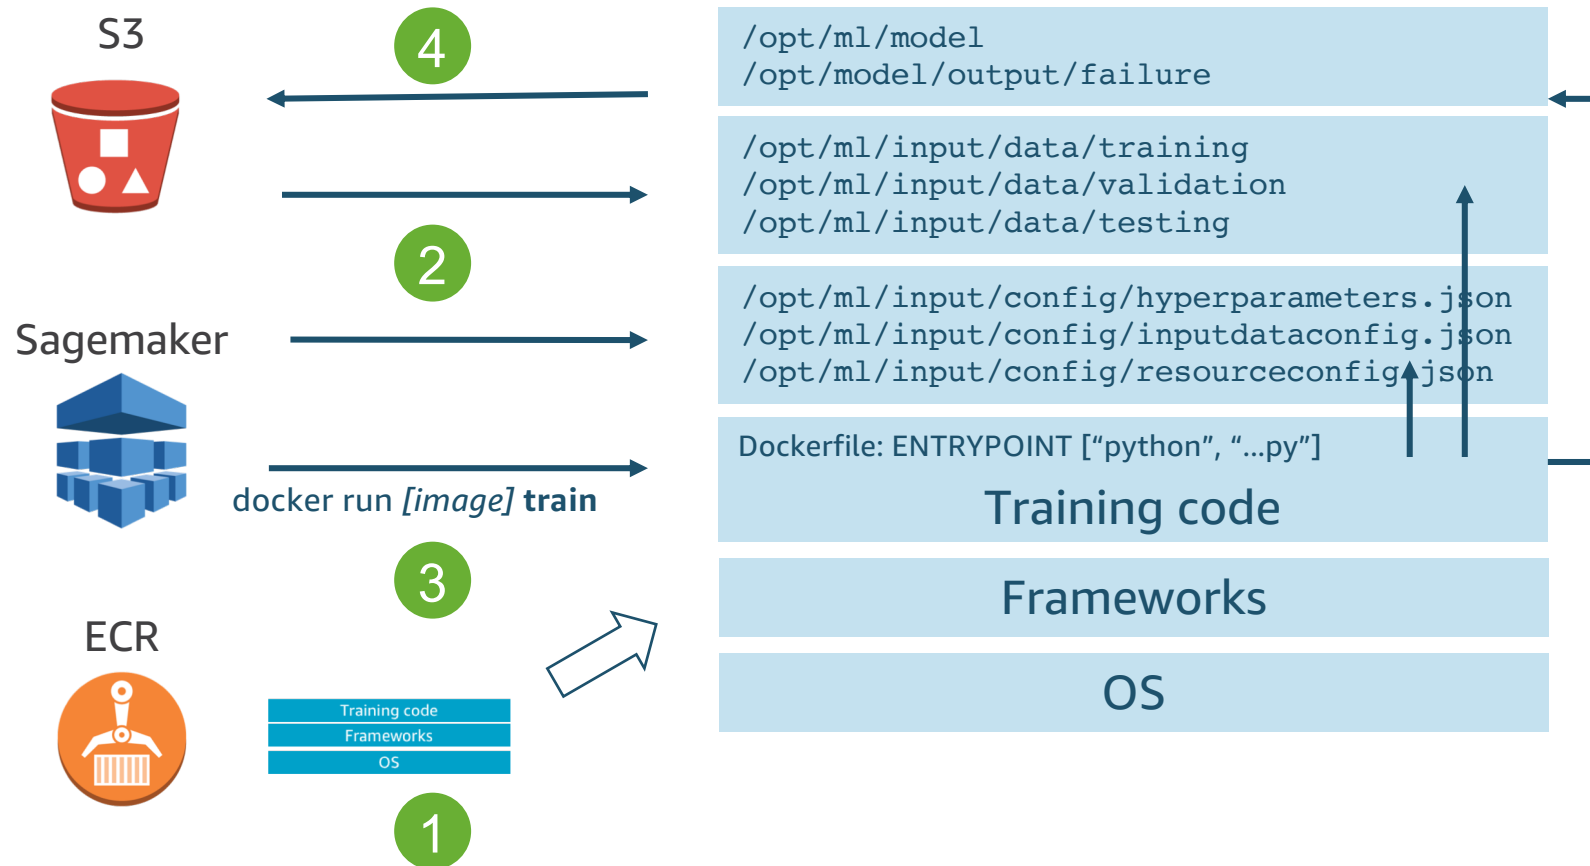


Amazon ECS

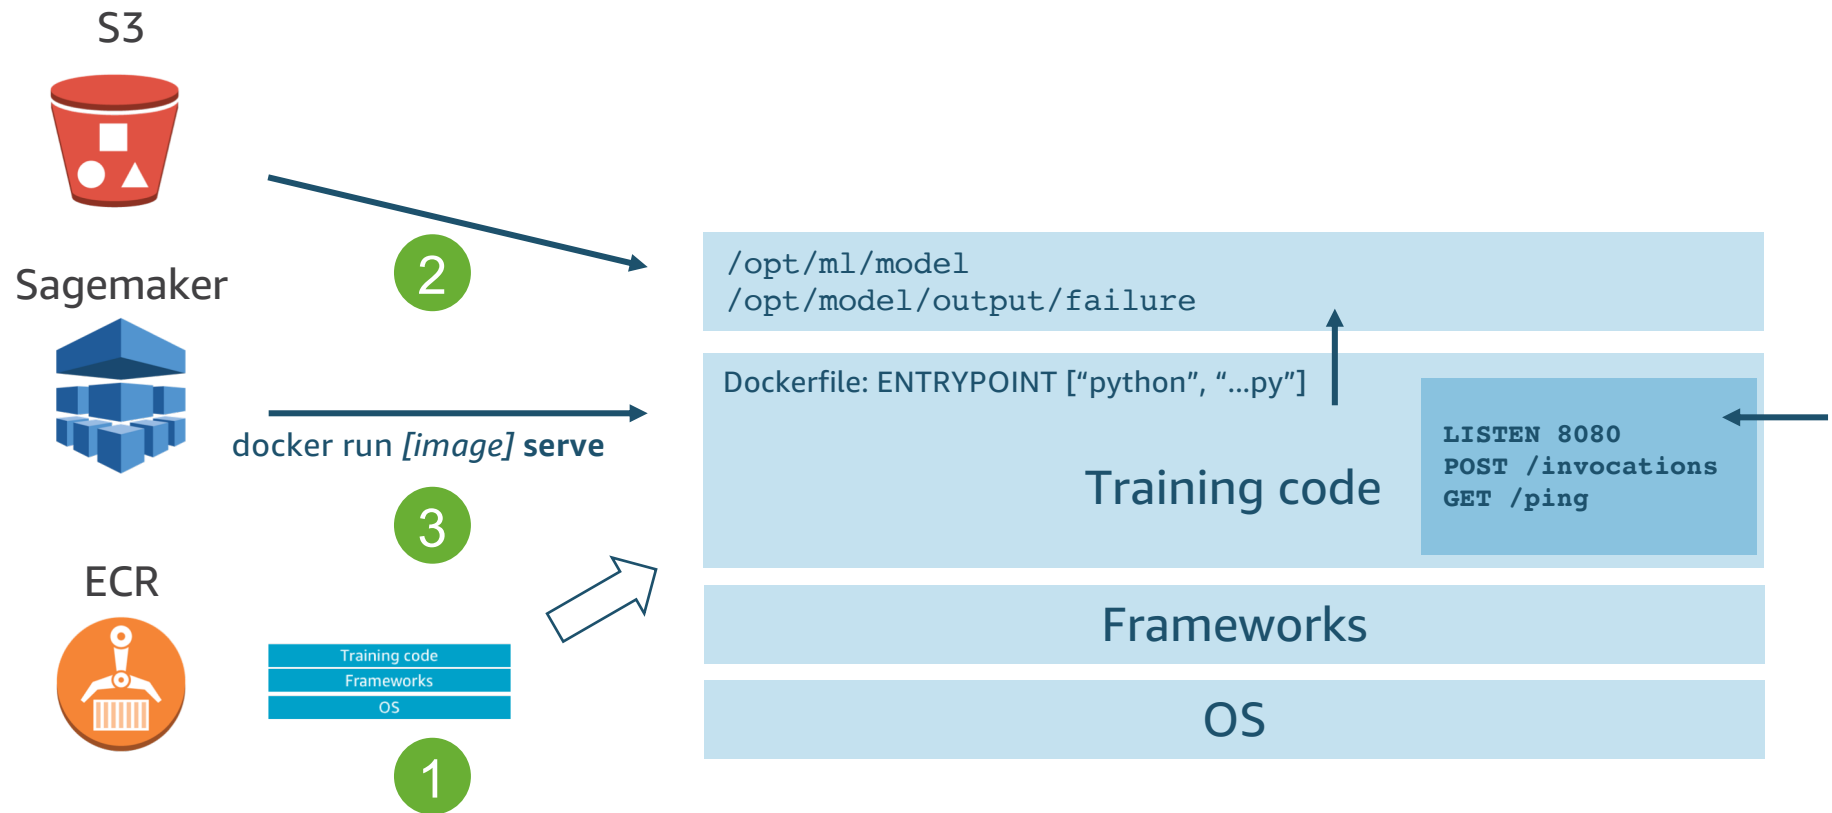
... publish to ECS



# BYOM - Training



# BYOM - Inference



# Questions